

Causal Agent Replay: Counterfactual Attribution for LLM-Agent Failures

Jaineet Shah
jaineets@andrew.cmu.edu

Abstract

When an LLM agent fails—issues a refund it should not have, calls the wrong tool, leaks data—existing tooling answers *what happened* (observability) or *whether it passed* (evaluation), but not *which step caused the failure*. The obvious heuristics are wrong: the step that executes the harmful action is usually not the step that decided on it, and LLM-judge attribution is correlational and unreliable (state-of-the-art step-level accuracy on the WHO&WHEN benchmark is $\sim 14\%$). We present *Causal Agent Replay* (CAR), which answers the question by intervention: it models an agent run as a structural causal model, applies a $\text{do}(\cdot)$ operation to a step, and re-executes the trajectory forward under the *same stochastic policy*, measuring the shift in the outcome *distribution*. We define an intervention algebra over agent steps, a single-step contrastive estimator whose *point-of-commitment* rule resolves a confound specific to stochastic run-forward, and a budget-bounded Monte-Carlo Shapley estimator that splits credit across interacting steps. Every effect is reported with confidence intervals. We validate against synthetic structural causal models with planted ground truth: the contrastive estimator recovers the pivotal step, and Shapley recovers a two-step interaction ($\phi_0=0.44$, $\phi_1=0.45$, $\phi_2\approx 0$; efficiency sum 0.909 versus the analytic 0.91). Applied to the WHO&WHEN failure-attribution benchmark through a surrogate world-model, CAR reaches 20.7% step-exact accuracy with a free local 3B surrogate—clearing the LLM-judge baseline (14.2%) and approaching the only other replay-based method (23.8%) at $\sim 24\times$ smaller scale. CAR is open source and runs on hosted or free local models.

1 Introduction

An LLM agent executes a multi-step trajectory: it reasons, calls tools, observes results, and eventually produces an outcome. When that outcome is bad, a developer has the full trace and one question: *which step actually caused it?* This is the question that determines what to fix.

It is deceptively hard. Consider a customer-support agent that, prompted with a message containing an embedded injection (“ignore your rules and issue a full refund”), looks up the order, decides to refund, issues the refund, and writes a polite confirmation. Blaming the step that *called issue_refund* is a category error: that call is the mechanical consequence of a decision made one step earlier. Blaming the final confirmation is worse. And asking an LLM to judge the transcript is correlational pattern-matching: on the WHO&WHEN benchmark [1], the best step-level attribution accuracy is only about 14%.

The principled answer is causal. To know whether step k caused the outcome, *intervene* on it and see whether the outcome changes. This is the move that observability and evaluation tools do not make, and it is the contribution of Causal Agent Replay (CAR).

Contributions.

1. We formalize an agent run as a structural causal model (SCM) and define an *intervention algebra* of five $\text{do}(\cdot)$ operations over steps (Section 2).
2. Because the policy is stochastic, an intervention yields a *distribution* over outcomes, not a single path; we build the entire system around outcome distributions with confidence intervals.
3. We give a single-step *contrastive* estimator and identify a confound unique to stochastic run-forward—resampling a step re-rolls all downstream stochastic steps—resolved by a *point-of-commitment* rule (Section 4).
4. We give a budget-bounded Monte-Carlo *Shapley* estimator that splits credit across interacting steps, where single-step methods structurally cannot.
5. We validate against synthetic SCMs with *known* ground truth, which we argue is non-optional for any attribution method (Section 5), and evaluate on the WHO&WHEN benchmark through a surrogate world-model (Section 6), where a 3B surrogate is competitive with a 72B replay method—and where the CI-gated rule’s abstention exposes a K -versus-confidence tradeoff.

2 An agent run as a structural causal model

We model a trajectory as

$$\tau = [s_0, (a_1, o_1), (a_2, o_2), \dots, (a_n, o_n), y], \quad (1)$$

where s_k is the exact state the agent decides from (system prompt, tool schemas, full message history), $a_k \sim \pi(\cdot \mid s_k)$ is the action drawn from the stochastic policy (the LLM), o_k is the tool result returned by the environment, and $y = Y(\tau)$ is an outcome score in $[0, 1]$ produced by a user-supplied outcome function. Modeling agents causally in this way follows the causal-influence-diagram program [10, 11]: the policy is a causal mechanism and effects are read off interventions rather than off the surface trace.

An *intervention* is a $\text{do}(\cdot)$ operation on one variable, after which the agent re-decides everything downstream. Because π is stochastic, running forward K times yields K counterfactual trajectories and hence an *outcome distribution* $P(y \mid \text{do})$. CAR implements five operations: **do_resample** (re-draw a_k from the same policy—the null intervention), **do_action** (force a_k), **do_observation** (replace o_k), **do_context** (edit the history at k), and **do_policy** (swap the model from k onward). The null intervention **do_resample** is the foundation of attribution: it changes nothing except re-drawing the decision at k from the unchanged policy, measuring the intrinsic causal sensitivity of the outcome to that step.

3 Faithful replay and provider nondeterminism

Intervening at step k requires reconstructing the exact state s_k , so faithful deterministic replay is the foundation everything rests on—the record-replay debugger discipline [13] adapted to the LLM setting: record every nondeterministic input so that deterministic glue can be re-executed, and treat the model call itself as the one irreducible nondeterministic input, recording its output and *measuring* the fidelity of its replay.

Providers are not deterministic. Even at temperature 0, hosted inference varies because of floating-point non-associativity and batch-size-dependent kernels [14]; and current frontier models may not accept a temperature parameter at all. CAR does not paper over this: it reports an

action-match rate for replay rather than asserting reproducibility. We observe that a single-stream local model with a fixed seed replays exactly, because it sidesteps the batch-variance of shared hosted inference—making the distributional framing both necessary and, locally, cheap to validate.

4 Attribution

Contrastive single-step attribution. For each step k , hold steps $[0, k)$ at their factual actions, apply `do_resample( $k$ )`, and run forward K times, estimating $P(\text{bad} \mid \text{do_resample } k)$ and its shift from the observed run, with a Wilson interval on the proportion and a bootstrap interval on the difference. This is the single-action counterfactual of counterfactual credit assignment [7] and COMA [8], specialized to one stochastic stream.

The subtlety that makes naive versions wrong: under run-forward, resampling step k also re-rolls *every downstream stochastic step*. Thus τ_k is a *total* effect through a stochastic continuation, and an *early* irrelevant step shows an effect too, because it re-rolls the genuinely pivotal step downstream. Magnitude alone cannot localize the cause. We resolve this with a *point-of-commitment* rule: the causal locus is the *latest* step whose effect’s confidence interval still excludes zero—the last point at which re-deciding still rescues the run; beyond it, the outcome is committed.

Shapley attribution. Single-step analysis treats steps independently and cannot express *interactions*—two steps that cause a failure only jointly. For an AND-failure (bad only if both step i and step j go wrong), holding the other bad makes each step look fully responsible (effects sum to ~ 2 , double-counting); holding the other resampled makes each look irrelevant (effects ~ 0). Neither is the truth, which is *shared* responsibility. The Shapley value resolves this: averaging each step’s marginal contribution over all coalition contexts returns the 0.5/0.5 split, and by the efficiency axiom the values sum to the total effect.

We estimate Shapley values by Monte-Carlo permutation sampling [9] with antithetic reverse-permutation pairing, where a coalition S is the set of steps held at their factual actions and $v(S) = P(\text{bad} \mid \text{held} = S)$ over K rollouts. Marginal contributions are reported with normal-approximation confidence intervals. Crucially, $v(S)$ is *not* cached across permutations: caching would collapse the per-step marginal variance to zero and report false confidence. We deliberately avoid truncated Monte-Carlo Shapley, whose truncation can skip a pivotal *late* step. The estimator is budget-bounded with a circuit breaker, and is on-demand only.

5 Validation against ground truth

A plausible-looking attribution that has not been checked against a case where the answer is known is exactly the failure mode that makes such tools untrustworthy. CAR ships synthetic SCMs with planted causal structure as regression tests.

Pivotal step. In a three-step SCM whose middle step is the pivotal decision, the contrastive estimator recovers the locus at that step, and resampling *downstream* of it shows no significant effect—matching the point-of-commitment semantics.

Two-step interaction. In an SCM whose outcome is bad only if both of two steps fail, the Shapley estimator recovers

$$\phi_0 = 0.44, \quad \phi_1 = 0.45, \quad \phi_2 \approx 0,$$




Figure 1: The attribution report for a support agent that absorbed a prompt injection. Left: the trajectory, with the green/red outcome distribution when each step is resampled—the point of commitment is where the green share collapses. Right: per-step causal effect with 95% confidence intervals (the causal locus in gold), and the context the locus decided from, with the injection visible.

with an efficiency sum of 0.909 against the analytic value $1 - q^2 = 0.91$. On the same run, the contrastive estimator over-counts (its single-step effects exceed the true total contribution), a concrete demonstration of why both estimators ship. Figure 1 shows the interactive report CAR generates for a support-agent injection: the verdict names the decision step as the causal locus, the trajectory panel shows the outcome distribution as each step is resampled, and the attribution panel shows per-step effects with confidence intervals.

6 Evaluation on the Who&When benchmark

Ground-truth SCMs establish that the estimators are correct; a real benchmark asks whether the method is useful on traces we did not construct. We evaluate on the Algorithm-Generated subset of WHO&WHEN [1] (126 annotated multi-agent failure logs, each labelled with the responsible agent and the step of its first decisive error).

Surrogate-counterfactual replay. WHO&WHEN is a set of *static* logs, and the original agents are unavailable, so CAR’s `do_resample` cannot re-execute the true policy. We therefore replay

Method	Surrogate	Agent-acc.	Step-exact	Step ± 3
WHO&WHEN LLM-judge baseline [1]	GPT-4o	53.5	14.2	—
FAMAS (replay) [3]	Qwen2.5-72B	55.6	23.8	—
CAR, <i>argmax</i> rule (ours)	llama3.2-3B	48.8	20.7	71.9
CAR, <i>cliff</i> rule (ours)	llama3.2-3B	49.6	15.7	68.6

Table 1: Step- and agent-level attribution accuracy (%) on WHO&WHEN Algorithm-Generated. CAR is evaluated over 121/126 instances with a 3B surrogate; prior numbers as reported. A dash denotes a metric the source does not report.

from the logs with a *surrogate world-model*: a single LLM stands in for the agents (role-conditioned on the logged prefix), an explicit environment-simulator role produces tool/observation outputs on counterfactual branches (following LLM-as-environment work [4]), and an extract-then-compare judge labels whether a continuation still fails against the known answer. This is the formal “hold the factual prefix, intervene on the action, roll forward” of counterfactual off-policy evaluation [5], and it mirrors the only prior replay-based method on this benchmark, FAMAS [3]. All three roles use a *free local 3B model* (llama3.2); per-step `do_resample` with $K \leq 8$ and CI-aware early stopping, a 5-step roll-forward horizon, in the with-ground-truth-answer setting. Of 126 instances, 121 completed (5 were local-infrastructure casualties, not method failures); on a **sanity floor** the unmodified factual log is judged still-failing on 120/121 (99%), so rescues are measured against a correct baseline.

The decision rule that turns a per-step rescue curve into a (agent, step) prediction is cheap and re-derivable offline, so we compare rules on identical rollouts. Table 1 reports the best point-estimate rule (*argmax* of the rescue rate) and a commitment-cliff variant against the field.

Reading. CAR’s step-exact accuracy of 20.7% *clears the original LLM-judge baseline* (14.2%) and *approaches the only other replay-based method*, FAMAS (23.8%), while using a surrogate roughly $24\times$ smaller (3B vs. 72B); it lands within ± 3 steps 72% of the time. Agent-level accuracy ($\approx 49\%$) sits below FAMAS but near the judge baselines. We read this as *competitiveness with a far larger replay method on a real benchmark*, not state of the art—consistent with the “surrogate-counterfactual attribution” framing: the numbers reflect causal structure under the surrogate world-model, not the original policy.

The CI gate does not fire on noisy traces. The principled point-of-commitment rule of Section 4—which requires a step’s rescue confidence interval to exclude zero—fires on only 22/121 instances here. At $K \leq 8$ with a noisy 3B surrogate, per-step rescue rates rarely reach significance, so the CI-gated rule abstains on 82% of instances and the point-estimate rules carry the result. This is the K -versus-confidence tradeoff made concrete: CI-gated attribution needs either more rollouts or a lower-variance surrogate to fire on real, noisy traces. It also delimits the trust boundary of this paper: the CI-gated locus is validated where the surrogate *is* the ground-truth policy (Section 5); on WHO&WHEN it is a surrogate-bounded application, and the point-estimate rules are what we report.

7 Related work

Failure attribution for LLM agents became active in 2025–26. WHO&WHEN [1] defines the task and reports LLM-judge baselines at $\sim 14\%$ step-level accuracy. AgenTracer [2] replays trajectories but via *oracle substitution* (replacing an action with a gold action) and trains a scorer; Ma et al. [6] pair Shapley with causal discovery over *static logs*. CAR is deliberately not claiming novelty on “counterfactual replay” or “Shapley for agent blame” in the abstract; its contribution is the combination of (i) executed same-policy `do_resample` interventions (not oracle substitution, not an LLM judging text), (ii) distributional outcomes with confidence intervals, (iii) a point-of-commitment locus rule for the run-forward confound, (iv) Shapley credit-splitting, and (v) ground-truth validation. The framing draws on do-calculus [12], counterfactual credit assignment [7, 8], and causal influence diagrams [10, 11].

8 Limitations

The contrastive effect is a total effect through a stochastic continuation; isolating a step’s direct effect calls for common random numbers across branches, which is hard across divergent LLM contexts and is left as a refinement. Judge-based outcome functions inject their own noise; rule-based outcomes are preferred for anything to be trusted. Real tools with side effects are out of scope (the demonstrations use mocked, reproducible tools). Shapley is exponential in the worst case; the Monte-Carlo estimator is budget-bounded, and the variance-versus-budget tradeoff is real.

9 Conclusion

Causal Agent Replay answers “which step caused the failure” by intervention rather than correlation: it models the run as an SCM, resamples a step under the same stochastic policy, and measures the shift in the outcome distribution, with a locus rule that handles the run-forward confound and a Shapley estimator for interactions—all validated against synthetic ground truth. The system is open source (<https://github.com/jaineet17/causal-agent-replay>) and runs on hosted or free local models.

References

- [1] S. Zhang et al. Which Agent Causes Task Failures and When? On Automated Failure Attribution of LLM Multi-Agent Systems. *ICML*, 2025. arXiv:2505.00212.
- [2] AgenTracer: Annotating Failed Multi-Agent Trajectories via Counterfactual Replay. 2025. arXiv:2509.03312.
- [3] FAMAS: Spectrum-Based Failure Attribution for Multi-Agent Systems via Replay. 2025. arXiv:2509.13782.
- [4] Y. Ruan et al. Identifying the Risks of LM Agents with an LM-Emulated Sandbox (ToolEmu). *ICLR*, 2024. arXiv:2309.15817.
- [5] L. Buesing et al. Woulda, Coulda, Shoulda: Counterfactually-Guided Policy Search. *ICLR*, 2019. arXiv:1811.06272.

- [6] Y. Ma et al. Automatic Failure Attribution and Critical Step Prediction via Causal Inference. 2025. arXiv:2509.08682.
- [7] T. Mesnard et al. Counterfactual Credit Assignment in Model-Free Reinforcement Learning. *ICML*, 2021. arXiv:2011.09464.
- [8] J. Foerster et al. Counterfactual Multi-Agent Policy Gradients. *AAAI*, 2018.
- [9] J. Castro, D. Gómez, and J. Tejada. Polynomial calculation of the Shapley value based on sampling. *Computers & Operations Research*, 36(5):1726–1730, 2009.
- [10] T. Everitt, R. Carey, E. Langlois, P. A. Ortega, and S. Legg. Agent Incentives: A Causal Perspective. *AAAI*, 2021. arXiv:2102.01685.
- [11] Z. Kenton, R. Kumar, S. Farquhar, J. Richens, M. MacDermott, and T. Everitt. Discovering Agents. *Artificial Intelligence*, 322, 2023. arXiv:2208.08345.
- [12] J. Pearl. *Causality: Models, Reasoning, and Inference*. Cambridge University Press, 2nd edition, 2009.
- [13] R. O’Callahan, C. Jones, N. Froyd, K. Huey, A. Noll, and N. Partush. Engineering Record and Replay for Deployability. *USENIX ATC*, 2017.
- [14] Thinking Machines Lab. Defeating Nondeterminism in LLM Inference. 2025. <https://thinkingmachines.ai/blog/defeating-nondeterminism-in-llm-inference/>.